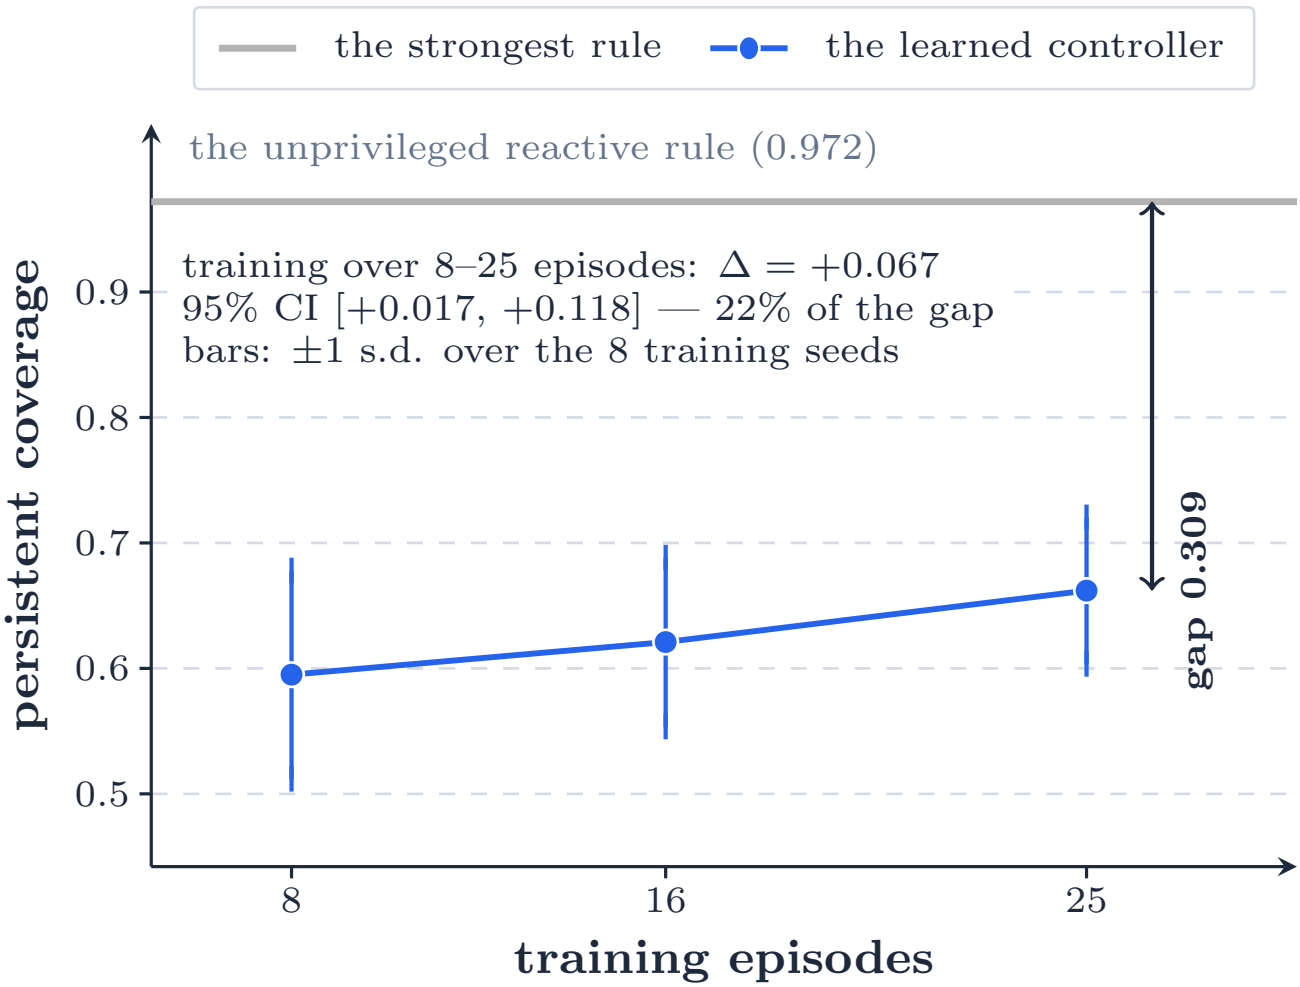
the strongest rule
the learned controller
the unprivileged reactive rule (0.972)
training over 8–25 episodes: Δ = +0.067
95% CI [+0.017, +0.118] — 22% of the gap
bars: ±1 s.d. over the 8 training seeds
gap 0.309
persistent coverage
0.9
0.8
0.7
0.6
0.5
8
16
25
training episodes