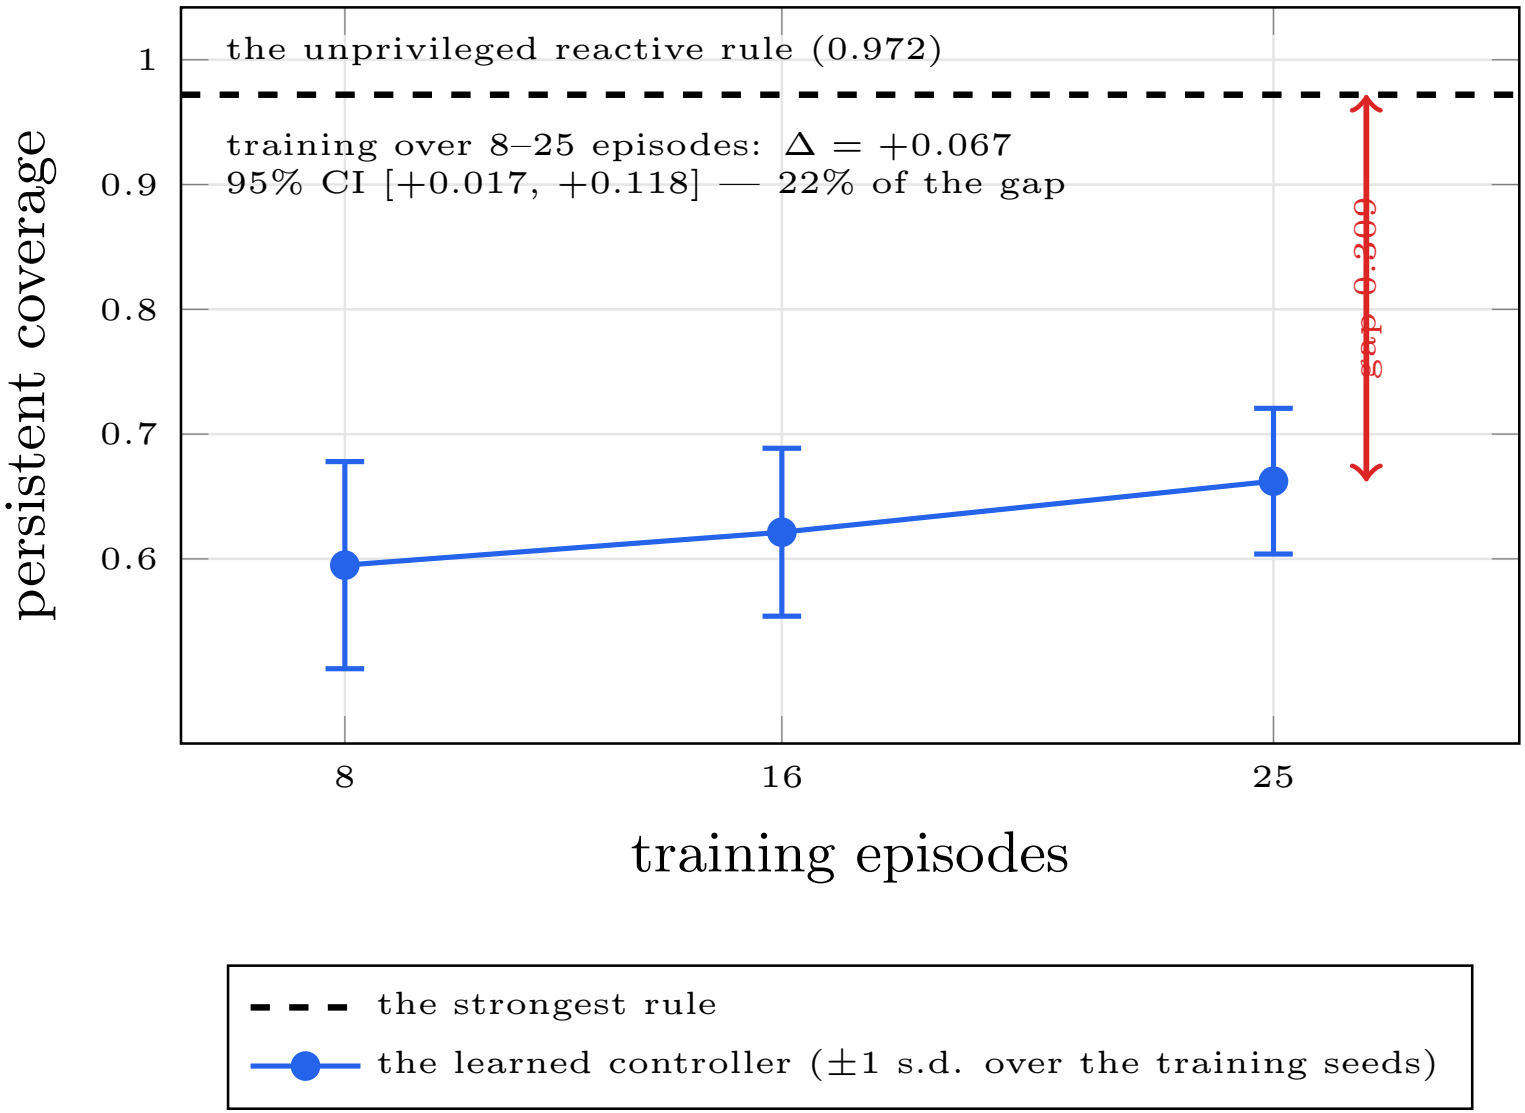
the unprivileged reactive rule (0.972)
training over 8–25 episodes: Δ = +0.067
95% CI [+0.017, +0.118] — 22% of the gap
gap 0.300
persistent coverage
1
0.9
0.8
0.7
0.6
8
16
25
training episodes
the strongest rule
the learned controller (±1 s.d. over the training seeds)