Persistent UAV-IoT Coverage: Project Report

The research log behind the paper: the protocol, the defects, the audits, the measurements, and what is still open

Generated from the artefacts by `analysis/make_report.py` — every number below is read from the file it was measured in.

1 The contract and the current state

The work runs to 26 October 2026. The deliverable is a paper whose every number can be traced to the artefact it was measured in, and whose comparative claims are derived in code rather than asserted, so that the paper says the right thing whichever way a measurement comes out. This report is the log of that work: the protocol, the defects that were found and fixed, the audits that were run and what they changed, the measurements, and the questions that are still open.

Gap to the strongest rule	**0.276** (95% interval [0.242, 0.309]), against the unprivileged reactive rule at 0.972.
The learned controller	**0.696** ± 0.027 over 5 training seeds.
Claim audit	35 items: 31 verified, 0 failed, 4 recorded for transparency.
Test suite	115 cases, all passing.
Device	every artefact records the device it was measured on; the device and the thread count are part of the protocol.

2 The measurement protocol

Every number in the paper and in this report is produced under one protocol:

1. **Training**: 25 episodes of 2000 control steps per training run. The base comparison trains five runs; the component ablation and the arms train three.
2. **Evaluation**: five scenes per policy, and the number of evaluation scenes is part of the protocol — a run with three is a different measurement, because a per-seed row is the mean over the scenes.
3. **The device and the thread count** are part of the protocol. The floating-point kernels differ between the CPU and the GPU and the order of the reductions differs with the thread count, so two runs of the same configuration that land on different devices are not the same experiment. The device is pinned, the thread count is pinned in code, and every artefact records the device.
4. **The evaluation is seeded** from the evaluation seed alone, so every controller is measured on the same random stream on the same scene — common random numbers — and a comparison between controllers is a comparison between policies rather than between the luck of the draw.
5. **The ablations are paired**: the variants share the training seeds, so a difference is measured on the same training runs and reported as a confidence interval on the paired difference.
6. **Inference is by interval, never by threshold**. An explanation is excluded only if the arm that tests it has an upper confidence limit below the gap; a null result is a statement about the resolution of that arm, and the resolution is reported with it.

3 What the measurements say

Controller	CK	AG	SC	Cov. deficit
Fixed	0.372	1.000	0.000	0.628
Greedy	0.786	0.964	0.066	0.214
Reactive	0.953	1.000	0.751	0.047
Reactive (local assignment)	0.972	1.000	0.777	0.028
MAPPO (proposed)	0.696 ± 0.027	0.985	0.015	0.304
MAPPO (argmax)	0.411 ± 0.075	0.988	0.000	0.589
Oracle (ceiling)	1.000	1.000	1.000	0.000

The arms' verdict, derived from the intervals: Controlled arms bound how much of the gap each alternative explanation can account for: adding the agent identity (upper limit 0.377, margin -0.101); the full observation of the proposed configuration (upper limit 0.164, margin +0.112); removing the energy constraint (upper limit 0.200, margin +0.076); removing the threshold term (upper limit 0.066, margin +0.210); the previous energy rule (upper limit -0.003, margin +0.279), against a gap of 0.276. The arms whose upper limit reaches past the gap – adding the agent identity – are NOT excluded, and remain candidate explanations; the others are excluded as the whole of the explanation.

The training-budget sweep: 8 episodes: 0.609, 16 episodes: 0.640, 25 episodes: 0.698. the coverage moves by +0.089 over the range the sweep covers, which is about 32\{}% of the gap of 0.276, and the increments between successive budgets are +0.031, +0.058; but the last of them, +0.058, is only 1.2 standard errors from zero (the standard error of the difference is 0.048, at three training seeds per budget), so the sweep does not resolve what the budget contributes and the budget is neither excluded nor shown to be a cause of the shortfall.

4 The defects that were found and fixed

Every one of these was a defect in the work, not in the writing. The list is the reason the claim audit exists.

#	**Defect**	**Fix**
1	Simulator was static – no UAV ever moved	`13cb839`
2	Baselines accepted `seed` and ignored it	`_seeded()`; `d113e5c`
3	Time limit reported as termination -¿ GAE bootstrapped 0	`d113e5c`
4	Centred-ring start pre-solved the task	`957ee1c`
5	`run_scenarios` passed `seed=0` over the config's own seed	`957ee1c`
6	Reward eq (23) negated only U_t	`reward_from_terms`
7	Initial placement disconnected on 11/20 seeds	`_connected_initial_positions`
8	Backhaul modelled as co-channel interference	SNR-based backhaul
9	`snr_to_rate` off-by-one + floor made the PHY gate vacuous	fixed
10	Path-loss exponent additive -¿ LoS/NLoS erased	`pl_d0 + 10n log10(d/d0)`
11	Noise units: dBm treated as watts (30 dB error)	divide by 1000
12	Shadowing from the global RNG -¿ non-reciprocal	seeded per-episode table

the defect log, continued

#	Defect	Fix
13	`local_obs` neighbour stride 2 -¿ slots overlapped	concatenation + assertion
14	`STATE_DIM = 80` hardcoded (valid only for N=5)	`state_dim(cfg)`
15	GAE bootstrap at truncation used 0 instead of V(s_T+1)	fixed
16	Observation carried no agent identity	`use_agent_id`; `44206f5`
17	Observation could not express a return (no direction to a pad)	`use_pad_bearing`; `de3e1e6`
18	No indication of where the coverage deficit is	`map_grid`; `44206f5`
19	The energy state machine did not exist	env makes the transition; `de3e1e6`
20	A UAV with an empty pack still counted as serving	out of service; `de3e1e6`
21	Charging was uneconomic (billed the hover power)	rotors idle on the pad
22	The battery never bound at 400 steps	episode 2000 steps
23	The returns were formed from an untruncated advantage array	truncate before forming returns
24	Two graph computations per candidate action (13.7 ms/step)	shared prospective graph
25	MAPPO's value normalisation missing although asserted	`c4a71f6`
26	The comparison table named MARL baselines never evaluated	`baselines_rl.py`; "Evaluated" column
27	The scenario matrix crashed on N=3: the one-hot identity has length N and the global state carries the sensor positions and the pad occupancy, so the networks' input sizes depend on the configuration	one policy per distinct `(N, M, pads)`
28	THE EVALUATION OF A SAMPLED POLICY WAS AN UNCONTROLLED RANDOM DRAW. Eq. (13) specifies that the deployed policy samples, and `mappo.py` drew from the global torch RNG, which `eval_agent` never seeded. The same policy file evaluated three times gave CK 0.734, 0.701 and 0.671, with the same standard deviation to four decimals. The paper's numbers could not be reproduced from the stored artefact, and controller-to-controller comparisons carried an unquantified sampling component	`uav_wireless/seeding.py`; 5 new tests; audit item R1
29	The component ablation trained ONE seed per variant at 10 episodes, while the base comparison trained five seeds at 25 – so its deltas mixed the component with the budget and its spread was unmeasurable. It also disagreed with the observation arms (which said the full configuration was no better than the bare one)	`run_ablation(..., seeds=)`, budget aligned to the base comparison, paired differences, incremental writes
30	The main table's learned row was ONE policy while the text claimed five training seeds	the row is the mean over the five runs with the seed-to-seed spread; every run's policy is saved

the defect log, continued

#	Defect	Fix
31	The assembler's "undefined references" check matched the LaTeX FONT warnings; the scenario table's variable names were set in text mode, so the table did not compile	precise check; math mode; `tabcolsep`; `raggedbottom`
32	The ablation's "is this component measurable" rule was not a test. It compared a component's difference with the learner's own seed-to-seed spread, a rule with a per-arm null error rate of about a third at three training seeds, and it bought no power	paired t intervals on the per-seed differences, plus a resolution statement; the interpretive sentences are DERIVED in the token filler, not asserted in the prose
33	The ablation runner wrote its incremental log with `dict(row, protocol)`, which raised on a duplicated key and killed the stage after two training runs	the payload is built by UPDATE; the runner RESUMES (a row measured under the same protocol is reused) and the CSV is rewritten after every row
34	The base comparison's CSV took its fieldnames from the first row, so the learned rows' extra columns raised on write	the fieldnames are the union over the rows; the base stage gained a `--resume` path that evaluates the stored policies instead of retraining them
35	The ablation's "is this component measurable" rule was not a test (audit 4)	paired t intervals on the per-seed differences, and a resolution statement; the interpretive sentences are derived in the token filler
36	The paper's central exclusion argument was overstated (audit 3)	the abstract, the introduction and the conclusion now state the descriptive gap, the interval, the resolution, and the causes the arms do NOT exclude

5 The audits, and what they changed

Audit 1 (done): the seven load-bearing claims. Found 4 wrong as stated; all corrected. Audit 2 (done): the reward terms, the mask and the joint projection, the H-step lookahead, the energy budget, the value normalisation, the metrics, the oracle normalisation, the evaluation protocol. 25 items, 21 verified, 0 failed, 4 informational. Audit 3 (DONE, 20 Oct 23:15): the model and the results' claims. Brief: \exttt{C:\sers\bust\ppData\ocal\emp\pencode\udit3_brief.txt}\exttt (9 items). It found, and these are now fixed: * The central exclusion argument was OVERSTATED. "What remains is credit assignment" is not supported by three non-detections at three training seeds: a non-detection is a statement about the arm's resolution, and a sufficiency test cannot eliminate a cause that is necessary but not sufficient. The abstract, the introduction and the conclusion now say the residual is CONSISTENT WITH the credit assignment and with several other causes the arms do not test, and that the paper does not claim to have excluded them. * The "loses on its own objective therefore the learning pipeline is at fault" inference was a false dichotomy -¿ fixed. * The constrained ceiling IS 1.000, because the oracle's construction is itself admissible under the safety filter, so the paper's claim that an informative constrained ceiling was missed was false -¿ fixed. * The hop rule constrains the MOVER's eccentricity, not the graph's diameter, and can freeze an agent in a dispersed graph -¿ stated. * The service-continuity gap separates the two controllers far more sharply than the coverage gap -¿ now reported beside it. * The spread quoted for a learned controller contains the rollout-to-rollout variation, and the cancellation of the scene component in a difference is an ASSUMPTION -¿ both stated. * The framework's contribution needed a measurement. On the report's own infeasible N=5 configuration the reactive heuristic still reaches CK 0.648 with full connectivity while the joint metric is 0.000 and the oracle reaches 0.990: neither metric alone distinguishes the instance from a feasible one, and it is the certificate that makes the reading decidable. Measured

by \extttanalysis/infeasible_instance.py\exttt, audit item I1. * The rule's privileged access was an untested alternative explanation. Now measured: \extttanalysis/infeasible_instance.py\exttt aside, the arm \extttReactive (local assignment)\exttt in \exttttrain.py\exttt runs the same rule without the global coverage state, and the verdict is derived from the measurement in the token filler. Audit 4 (DONE, 20 Oct 22:10): the statistics of the re-measurement protocol itself. Brief: \extttC:\sers\bust\ppData\ocal\emp\pencode\udit4_brief.txt\exttt. It found the following, and all of it is now fixed: * The "yardstick" rule was not a test. Comparing a component's difference with the learner's own seed-to-seed spread has a per-arm null error rate of about a third at three training seeds, and it buys no power: the measurement cannot detect a true effect of the size it flags. -¿ replaced by paired t intervals on the per-seed differences, and the paper says "no effect larger than X was DETECTED at this budget and this seed count" rather than "the component is inert". * The negative conclusion's mechanistic half was overstated. The arms are null results used as evidence of absence. -¿ the abstract, the introduction and the conclusion now state the descriptive gap with its interval, and list the claims the measurements do NOT support. * The quoted uncertainty omitted the scene variance. -¿ the variance is now decomposed: the seed component is the uncertainty of a comparison, the scene component is common to all controllers and cancels in every difference. * The matched budget does not control for budget-conditionality or for redundancy between components. -¿ the paper says so, and says a budget sweep would be needed. * "A training run is deterministic given its seed" was not established. -¿ now established by two tests and by audit item T1 (the seed is set BEFORE the networks are constructed, so the initialisation is seeded too). * The deterministic-vs-sampled comparison. The paper's argmax row already settles the sampling-handicap objection, and the heuristic's information advantage is now stated as a limitation. * The audit's recommendations NOT taken, and why: raising the seed count to 10+ (the compute does not fit the deadline, so the paper reports intervals and a resolution instead of verdicts), decoupling the scene and action streams (would need the whole matrix re-measured), and a budget sweep (same reason). Each audit gets a SELF-CONTAINED brief: the setting, the model, the measured numbers, and the exact claims to check, because a fresh subagent has no context. It is instructed to work purely mathematically (no web, no code, no repository) and to be adversarial rather than agreeable.

6 Negative results, kept on the record

- The first OMNeT++ validation counted MAC frames and reported links closing at 50 km. The control (TX crushed 60 dB) returned the same count. A physical-layer claim without a negative control is not a measurement. The paper claims only model equivalence (0.012 dB). - The research found NO 2023-2026 paper that certifies a persistent-coverage instance INFEASIBLE, and NO paper that names the all-constraints-simultaneous metric as degenerate. Both are the paper's novelty claim. - Closest prior work: "Feasibility Ceilings Before Learning" (arXiv:2610.20823, Oct 2026). - The Reactive heuristic is a strong baseline (CK 0.953). It is NOT legitimate to hide it: it is reported in every table. - The reward's articulation-point term (lam_C) penalises the very spreading behaviour coverage requires. The ablation measures it.

7 The artefacts

Every number above is read from one of these files, and the paper is generated from the same ones.

results_main.csv	the main table: every controller, five training seeds for the learned one
results_ablation.csv	the component ablation, twelve variants
results_ablation_stress.csv	the ablation on the combined-stress scenario
results_scenarios.csv	the nine scenarios and their variants
analysis/ab_observation.json	the observation arms
analysis/ab_energy_management.json	the energy arms
analysis/ab_credit.json	the credit-assignment arm
analysis/ab_budget.json	the training-budget sweep
analysis/ab_map_resolution.json	the map-resolution sweep
analysis/claim_audit.json	the claim audit, item by item
analysis/connectivity_accounting.json	the connectivity accounting
analysis/corner_transfer.json	the corner certificate's transfer

8 Reproducing this report

```
cd ml_impl
python analysis/make_report.py     # this report
python analysis/assemble_results.py # the paper
python analysis/final_checks.py    # the gate
python analysis/claim_audit.py     # the audit
python tests/run_tests.py -q        # the tests
```